

OCR Canary Test Document

This document serves as a canary test artifact for the Industrial Adaptive OCR pipeline. It contains two pages of structured content designed to validate end-to-end processing through extraction, optical character recognition, assembly, and compression stages. The first page provides a block of plain English text to verify that the PaddleOCR engine can accurately detect and transcribe standard typeset characters at three hundred dots per inch. Successful recognition of this paragraph confirms that the GPU worker, language detection, and text-layer embedding stages are functioning correctly. The second page contains a tabular data sample that exercises the layout analysis and structured extraction components of the pipeline.

Sample Data Table

Name	Date	Amount
Alice Johnson	2026-01-15	\$1,250.00
Robert Chen	2026-02-03	\$3,780.50
Maria Garcia	2026-02-28	\$925.75